

Names: Bob Skowron, Jason Walker

Keys: rskowron, jwalker

SVN: jwalker: https://svn.seas.wustl.edu/repositories/jwalker/cse427s_f17/

1.
 - a. The *InputFormat* to use is *KeyValueTextInputFormat*. The default split is tab so no extra settings necessary.
Code to add to driver: *gob.setInputFormatClass(KeyValueTextInputFormat.class)*
To get the file path, we use the context to get the current input split and then read the path and name
FileSplit fileSplit = (FileSplit)context.getInputSplit();
String filename = fileSplit.getPath().getName();
 - b. See SVN
 - c. Mapper output for lines from Hamlet:
have hamlet@282
heaven hamlet@282
and hamlet@282
earth hamlet@282
there hamlet@133
are hamlet@133
more hamlet@133
things hamlet@133
in hamlet@133
heaven hamlet@133
and hamlet@133
earth hamlet@133

Names: Bob Skowron, Jason Walker

Keys: rskowron, jwalker

SVN: jwalker: https://svn.seas.wustl.edu/repositories/jwalker/cse427s_f17/

2. a. Cosine similarity: $\frac{r_x^T \cdot r_y}{\|r_x\| \cdot \|r_y\|}$

This is equivalent to: $\frac{\sum_{s \in S_{xy}} r_{xs} * r_{ys}}{\sqrt{\sum_{s \in S_{xy}} r_{xs}^2} * \sqrt{\sum_{s \in S_{xy}} r_{ys}^2}}$

Normalizing (i.e. subtracting the relevant means) gives: $\frac{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x) * (r_{ys} - \bar{r}_y)}{\sqrt{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)^2} * \sqrt{\sum_{s \in S_{xy}} (r_{ys} - \bar{r}_y)^2}}$

Which is the Pearson correlation.

- b.
 - An advantage of normalization is to remove bias. We can remove users who are overly critical (all low scores) or overly enthusiastic (all high rankings)
 - With Pearson correlation you need to compute and store the average rating for each user. This requires us to process the entire set of ratings.
- c. One disadvantage of the Jaccard similarity is that it does not take into account the value of the rating. Thus, even if two users rate the exact same items with opposite ratings it would have a high Jaccard similarity ranking. To overcome this problem, you could initially group by rating and then perform the similarity measure.

Names: Bob Skowron, Jason Walker

Keys: rskowron, jwalker

SVN: jwalker: https://svn.seas.wustl.edu/repositories/jwalker/cse427s_f17/

3. a. The dual approach can be more efficient if the $\#users \gg \#items$. Also, it can be easier to find similar items.

b.

User	Movie1	Movie2	Movie3	Movie5
user1	1	3	2	
user2		2	3	5
user3	1	2		

Mapper Output (pairs of movies and ratings):

((movie1,movie2), (1,3))

((movie1,movie3), (1,2))

((movie2,movie3), (3,2))

((movie2,movie3), (2,3))

((movie3,movie5), (3,5))

((movie1,movie2), (1,2))

Reducer Input:

((movie1,movie2), [(1,3),(1,2)])

((movie1,movie3), (1,2))

((movie2,movie3), [(3,2),(2,3)])

((movie3,movie5), (3,5))

((movie1,movie2), (1,2))

Reducer Output (similarity measure of movie pairs):

((movie1,movie2), .99)

((movie1,movie3), 0)

((movie1,movie5), 0)

((movie2,movie3), .92)

((movie2,movie5), 0)

((movie3,movie5), 0)