

Names: Bob Skowron, Jason Walker

Keys: rskowron, jwalker

SVN: jwalker: https://svn.seas.wustl.edu/repositories/jwalker/cse427s_f17/

1. a.
- b.
- c.

Names: Bob Skowron, Jason Walker

Keys: rskowron, jwalker

SVN: jwalker: https://svn.seas.wustl.edu/repositories/jwalker/cse427s_f17/

2.
 - a.
 - b.
 - c.
 - d.
 - e.
 - f.

Names: Bob Skowron, Jason Walker

Keys: rskowron, jwalker

SVN: jwalker: https://svn.seas.wustl.edu/repositories/jwalker/cse427s_f17/

3.
 - a. See SVN
 - b. *spark-submit CountJPGs.py /loudacre/weblogs*
Number of JPGs: 64,978
The Driver program and processing are done locally. The result is stored locally.
 - c. *spark-submit -master yarn-client CountJPGs.py /loudacre/weblogs*
The Driver program is run locally. The processing is done on the cluster. The result is stored in HDFS.
 - d. 1 stage and 311 tasks were executed
 - e. *spark-submit -master yarn-cluster CountJPGs.py /loudacre/weblogs*
The Driver program and processing are completed on the cluster. The result is stored in HDFS.

Names: Bob Skowron, Jason Walker

Keys: rskowron, jwalker

SVN: jwalker: https://svn.seas.wustl.edu/repositories/jwalker/cse427s_f17/

4.
 - a.
 - b.
 - c.
 - d.

Names: Bob Skowron, Jason Walker

Keys: rskowron, jwalker

SVN: jwalker: https://svn.seas.wustl.edu/repositories/jwalker/cse427s_f17/

5. a.
- b.